



This is a copy of a chat between Claude and sandy. Content may include unverified or unsafe content that do not represent the views of Anthropic. Shared snapshot may contain attachments and data not displayed here.

[Report](#)

10 preparing hiring screening test for apple email abuse annotation team hiring

Here's a 10-question screening test for the Apple Email Abuse Annotation role:



Log in to see interactive visuals

The test covers the key competencies Apple typically evaluates for an email abuse annotation role:

Categories tested:

- Classification (phishing vs. spam vs. malware)
- Severity rating and harm amplifiers (vulnerable populations, technical deception)
- Signal identification (SPF/DKIM, spoofed headers)
- Edge case and ambiguity judgment (joking intent, retroactive domain flagging)
- Policy knowledge (CAN-SPAM)
- Annotation consistency and drift detection
- Inter-annotator disagreement resolution
- Quality metrics (precision vs. recall trade-offs)

Scoring thresholds used: 8-10 = strong pass, 6-7 = borderline, below 6 = does not meet threshold. You can adjust these based on your team's calibration standards. Want me to add more questions, adjust difficulty, or export this as a shareable document?

we have separate test for category annotation

Got it! Let me update the test to remove the classification questions and focus purely on the other competencies.



Log in to see interactive visuals

Updated — classification questions removed. The 10 questions now cover:

- **Severity rating** (2 questions) — signal-based escalation, vulnerable population targeting
- **Edge case judgment** (2 questions) — intent vs. content, retroactive domain flagging
- **Inter-annotator agreement** (1 question) — severity disagreement resolution
- **Annotation consistency** (2 questions) — streak drift, intra-annotator variance
- **Policy knowledge** (2 questions) — CAN-SPAM, GDPR consent
- **Quality metrics** (1 question) — precision/recall trade-offs by context

Want to adjust any categories, add more questions to a specific area, or tune the difficulty?

make annotation scenario based



Log in to see interactive visuals

Fully scenario-based now. Each question puts the candidate inside a real annotation situation — they see the actual email (with headers and body), any annotator context notes, and then make the decision. The format across the 10 scenarios:

- 4 scenarios include a rendered email mockup (from/to/subject/body)
- 4 scenarios use annotator context notes (e.g. disagreement between two raters, a threat intel flag)
- 2 scenarios combine both

The categories stay the same — severity, edge case judgment, inter-annotator agreement, consistency, policy, and quality metrics — but every question now tests whether the candidate can apply judgment to a concrete situation rather than answer abstractly.

Want to adjust any scenario, add more edge cases, or make specific questions harder?

Make the test with scenarios where someone has to annotate it



Log in to see interactive visuals

Completely rebuilt as a hands-on annotation workspace. For each of the 10 emails, candidates must:

1. **Read the email** — full headers (From, Reply-To, To, Subject) and body, plus any annotator context notes
2. **Set severity** — Not abusive / Low / Medium / High / Critical
3. **Tag abuse signals** — pick all that apply from the signal taxonomy (spoofed sender, fake domain, urgency, impersonation, vulnerable target, etc.)
4. **Choose recommended action** — No action / Filter to junk / Remove & block / Escalate

Each scenario is scored out of 3 points (severity, signals, action) with partial credit for close calls. After all 10 are submitted, the results screen shows total score and a per-scenario breakdown.

The 10 emails cover the full range: obvious phishing, sophisticated Reply-To spoofing, legitimate Apple emails wrongly reported, spam vs phishing edge cases, a tech support scam targeting an elderly user, and a few genuinely ambiguous cases that test over-labelling instincts.

Want to add more scenarios, adjust the signal taxonomy, or change the scoring weights?

Can you give this in PDF format with answer key

Viewed a file, ran a command, created a file >

